


# Scalable Non-blocking Preconditioned Conjugate Gradient Methods

Paul Eller and William Gropp  
University of Illinois at Urbana-Champaign  
Department of Computer Science



Supercomputing 16



# Motivation

- To achieve best performance on extreme-scale systems, we need more scalable methods
- Blocking collectives are barrier to scalability
  - Increasing cost as core count increases
  - Requires synchronization across cores
- Noise causes performance variation across cores
- Methods with frequent synchronization will struggle to perform well as supercomputers grow larger




- What did we want to do?
  - Study preconditioned conjugate gradient (PCG) methods using non-blocking allreduces
  - Previous research answered many questions about numerical properties
  - But when do scalable PCG methods actually outperform standard PCG?
- What did we do?
  - Theoretically analyzed performance using performance models
  - Started developing software to analyze scalable algorithms
  - Found conditions where scalable PCG methods thrive
  - Developed a new method, PIPE2CG, using performance results to guide method development



# Preconditioned Conjugate Gradient Method (PCG)

- Popular iterative method for solving sparse symmetric positive definite linear systems  $Ax = b$
- Preconditioners needed in practice to improve convergence
- Blocking allreduce is barrier to scalability for PCG
- Need to minimize allreduce cost and avoid synchronization




- Scalable approaches:
  - Reduce communication latency by combining multiple allreduces
  - Overlap communication and computation using non-blocking allreduces
- Potential to:
  - Hide or avoid most of allreduce cost
  - Avoid synchronization cost
- Deriving new methods:
  - Use recurrence relations to rearrange order of key kernels
  - Equivalent to PCG in exact arithmetic
  - Increases initialization and vector operations costs



# Scalable PCG Methods


## Single Allreduce PCG

D'Azevedo et al 1993




## Non-blocking PCG

Gropp 2010



## Pipelined PCG

Ghysels et al 2014




# Further Pipelined PCG Methods

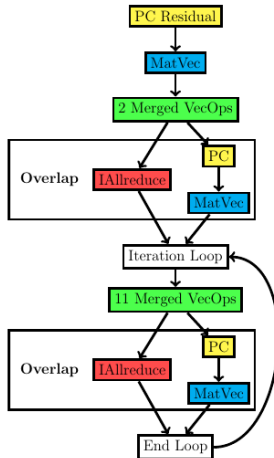
- PIPECG seems like reasonable starting point, but requires extra matvecs and PCs
- Instead start with three-term recurrence PCG
- PIPE2CG derived with similar process as two-term recurrence methods
- Three-term recurrences can produce less accurate residuals than two-term recurrences
- Slightly different approach will likely be needed to derive PIPE3CG, PIPE4CG, etc



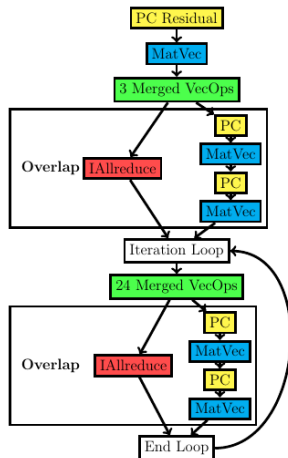
## Non-blocking PCG



## Pipelined PCG



## 2-Iteration Pipelined PCG





# Other Scalable PCG Methods

- Communication avoiding Krylov solvers:
  - Potential to reduce communication at multiple levels
  - Preconditioning matrix-powers kernel is difficult
  - PIPE2CG could use matrix-powers kernel to get non-blocking communication-avoiding solver
- Other Methods:
  - Pipelined GMRES: Similar approach as Pipelined PCG methods
  - Hierarchical and nested Krylov methods
  - Enlarged Krylov methods



# Merged Vector Operations

## Separate VecOps

```
AYPX(p, beta, u);  
AYPX(s, beta, w);  
Dot(p, s, gamma);
```



## Merged VecOps

```
for (i=0; i<n; i++) {  
    p[i]=u[i]+beta*p[i];  
    s[i]=w[i]+beta*s[i];  
    gamma+=p[i]*s[i];  
}
```

- Rearranging PCG introduces extra vector operations
- Minimize increased cost by performing vector operations element-wise
- Reduces vector reads
- Still requires extra vector writes




Figure: Vector reads and writes for 20 iterations of each method



# Non-blocking Allreduce

- Expected Usage:
  - Starts allreduce and returns control to user
  - Use wait to verify allreduce has completed
- Problem: May execute blocking allreduce in wait

Ideal Approach

```
MPI_Iallreduce();  
computation();  
MPI_Wait();
```

MPI\_Test Approach:

```
MPI_Iallreduce();  
for (...) {  
    chunk_comp();  
    MPI_Test();  
} MPI_Wait();
```

- Approaches:
  - MPI\_Test() - Break computation into chunks and give MPI control of thread
  - Progress threads - Dedicate one or more threads per node to communication
  - Hardware acceleration - Execute non-blocking allreduce in hardware



## ■ Experimental Setup

- Tests run on Blue Waters, a Cray XE6/XK7, with custom PETSc solvers
- Primarily use block-Jacobi incomplete-Cholesky preconditioner and Poisson matrices

## ■ Benchmarking Approach

- Use Hoefer et al 2015 and Hunold et al 2014 for guidance on producing accurate timings
- Developed code to collect parallel timings and compute statistics
- Produced iteration and 21-iteration test stats
- Primarily plot median runtimes, using 99% confidence intervals for 21-iteration tests



# Scalable PCG Performance Model

- Motivation:
  - Suboptimal non-blocking allreduce performance in practice
  - Performance variation at scale
- Model Setup:
  - LogGOPS model for parallel communication on Blue Waters
  - Netgauge for network parameters
  - Modified STREAM benchmark for computation parameters
- Overview of accuracy:
  - Accurately captures general solver trends
  - Improved cache, noise, and non-blocking allreduce models would improve runtime prediction accuracy



# Expected Performance

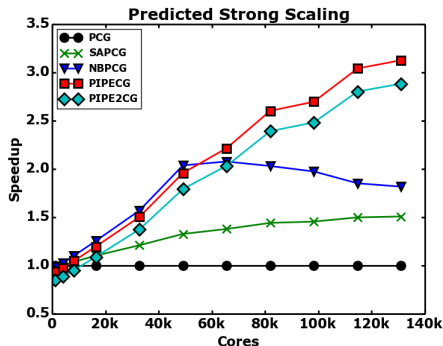


Figure: 27-point Poisson problem with  $512^3$  rows for 21 iteration tests

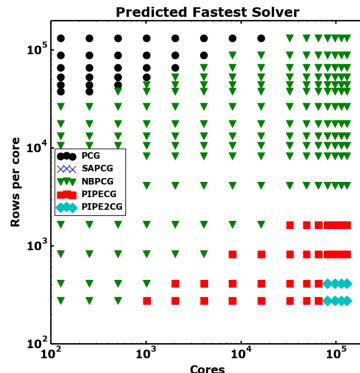


Figure: 27-point Poisson matrix for 21 iteration tests

# Weak Scaling Results




Figure: 27-point Poisson matrices with 4k rows per core

- Non-blocking solvers can hide communication costs and absorb noise to produce more consistent runtimes
- Increased initialization costs limit performance for non-blocking methods for 21-iteration tests

# Strong Scaling Results

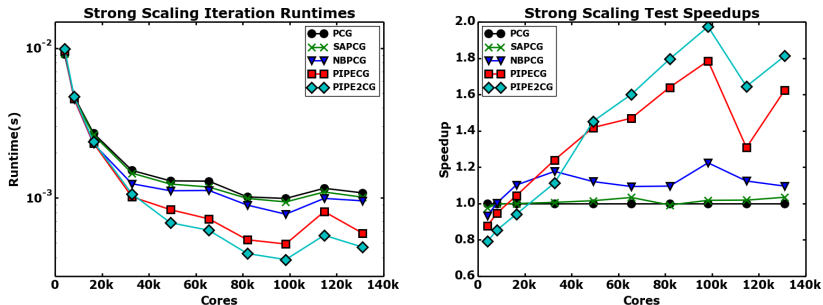



Figure: 27-point Poisson matrices with  $512^3$  rows


- Non-blocking methods struggle to scale further once computation cannot fully overlap communication
- Must overlap allreduce with more matrix kernels as work per core decreases and communication costs increase








(a) All cores



(b) Slowest core



(c) Prediction accuracy

Figure: Total runtime prediction accuracy using all cores and slowest core each iteration for 27-point Poisson matrices with 13k rows per core

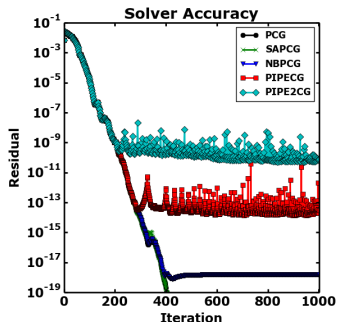


Figure: Computed residual norms for 5-point Poisson matrix with  $256^2$  rows

- Most tested matrices produced more accurate residuals
- Ghysels and Vanroose suggest residual replacement strategy
- Cools et al 2016 developed PIPECG method to monitor rounding error and use residual replacement as needed
- Effective for cheaper preconditioners and problems with  $O(1000)$  iteration counts

# Preconditioners




Figure: Strong scaling 27-point Poisson matrix with  $512^3$  rows




Figure: Weak scaling 3-d Laplacian matrix with 4k rows per core

- Assumed preconditioners call MPI routines enough to progress allreduce
- More detailed study needed to determine most scalable settings



# Weak Scaling Poisson Matrices




Figure: Iteration runtimes for matrices with about 220k nonzeros per process

- Matrices with more nonzeros per row can overlap allreduce with fewer rows per core, limiting overhead
- 2-d finite element matrices (9 and 18 nonzeros per row) show performance similar to 5-point Poisson matrices
- 3-d finite element matrices (80 nonzeros per row) show performance similar to 125-point Poisson matrices



- Scalable PCG methods:
  - More efficient communication
  - Increased initialization and vector operations costs
  - Effective for simpler preconditioners
  - Some speedup for more complex preconditioners without code modification
  - Hardware accelerated collectives may improve non-blocking allreduce



# How to use in practice

- Factors to consider for problems:
  - Expected iteration count
  - Matrix nonzeros per row
- Best Problems:
  - Enough iterations to converge to overcome increased initialization cost
  - More nonzeros per row in matrix



# How to use in practice

- Factors to consider for method:
  - Matrix/Vector rows per core
  - Core count
  - Expected system noise
- Method Selection:
  - Not a single fastest PCG method
  - Suite of scalable solvers provides best performance
  - Use least rearranged method capable of effectively overlapping allreduce with matrix kernels
  - May need extra rows per core in noisier environments



# Conclusions and Future Work

- Conclusions:
  - Can rearrange PCG to produce significant speedups at scale
  - Discovered conditions where scalable PCG methods outperform standard PCG using experiments on up to 128k cores of Blue Waters
  - Developed a new method, PIPE2CG, which outperforms other PCG methods for lower work per core and higher core counts
- Areas for Future Research:
  - Solver performance in noisy environments
  - Effectively using more complex preconditioners
  - Performance, accuracy, and usability within scientific applications
  - Non-blocking variations of other Krylov solvers
- Paper goes into much more detail on many topics





# Backup Slides

# Rearranging PCG

- Use recurrence relations to change vector used in key kernels
- Can then rearrange order of key kernels
- Example:
  - Start:  $r_j \leftarrow r_{j-1} - \alpha_j w_j$  and  $z_j \leftarrow Mr_j$
  - Replace:  $z_j \leftarrow M(r_{j-1} - \alpha_j w_j) \leftarrow Mr_{j-1} - \alpha_j Mw_j$
  - End:  $s_j \leftarrow Mw_j$  and  $z_j \leftarrow z_{j-1} - \alpha_j s_j$
- Equivalent to PCG in exact arithmetic
- Rearrangement produces additional initialization and vector operations costs



---

```

1:  $r_0 \leftarrow b - Ax_0$ 
2:  $z_0 \leftarrow Mr_0$ 

3:  $\gamma_0 \leftarrow (z_0, r_0)$  ▷ 2 Merged Operations
4:  $norm_0 \leftarrow \sqrt{(z_0, z_0)}$  ▷ Blocking Allreduce

5: for  $j=1,2,\dots$ , until convergence do
6:   if  $i$  then  $\beta_j \leftarrow \gamma_{j-1}/\gamma_{j-2}$ 
7:   else  $\beta_j \leftarrow 0.0$ 
8:    $p_j \leftarrow z_{j-1} + \beta p_{j-1}$ 
9:    $w_j \leftarrow Ap_j$ 
10:   $\delta_j \leftarrow (p_j, w_j)$  ▷ Blocking Allreduce

11:   $\alpha_j \leftarrow \gamma_{j-1}/\delta_j$ 
12:   $x_j \leftarrow x_{j-1} + \alpha_j p_j$ 
13:   $r_j \leftarrow r_{j-1} - \alpha_j w_j$ 
14:   $z_j \leftarrow Mr_j$ 

15:   $\gamma_j \leftarrow (z_j, r_j)$  ▷ 2 Merged Operations
16:   $norm_j \leftarrow \sqrt{(z_j, z_j)}$  ▷ Blocking Allreduce

```

---

- PCG contains four key kernels:
  - Matrix-vector multiply
  - Preconditioner
  - Vector operations
  - Allreduce
- PCG contains two blocking allreduces



# Single Allreduce PCG (SAPCG)

---

```
1:  $r_0 \leftarrow b - Ax_0$     $z_0 \leftarrow Mr_0$     $s_0 \leftarrow Az_0$ 
2:  $\delta_0 \leftarrow (z_0, s_0)$     $\gamma_0 \leftarrow (z_0, r_0)$ 
3:  $norm_0 \leftarrow \sqrt{(z_0, z_0)}$            ▷ 3 Merged Operations
4: Allreduce on  $\delta_0, \gamma_0, norm_0$        ▷ Blocking

5: for  $j=1,2,\dots$ ,until convergence do
6:   if  $j > 1$  then  $\beta_j \leftarrow \gamma_{j-1}/\gamma_{j-2}$ 
7:   else            $\beta_j \leftarrow 0.0$ 
8:    $p_j \leftarrow z_{j-1} + \beta_j p_{j-1}$ 
9:    $w_j \leftarrow s_{j-1} + \beta_j w_{j-1}$ 
10:   $\phi_j \leftarrow \delta_{j-1} - \beta_j^2 * \phi_{j-1} / \beta_{j-1}^2$ 

11:   $\alpha_j \leftarrow \beta_j / \phi_j$ 
12:   $x_j \leftarrow x_{j-1} + \alpha_j p_j$ 
13:   $r_j \leftarrow r_{j-1} - \alpha_j w_j$ 
14:   $z_j \leftarrow Mr_j$     $s_j \leftarrow Az_j$ 

15:   $\delta_j \leftarrow (z_j, s_j)$     $\gamma_j \leftarrow (z_j, r_j)$ 
16:   $norm_j \leftarrow \sqrt{(z_j, z_j)}$            ▷ 3 Merged Operations
17:  Allreduce on  $\delta_j, \gamma_j, norm_j$        ▷ Blocking
```

---

- Uses single blocking allreduce
- Algorithm from D'Azevedo et al 1993

# Non-blocking PCG (NBPCG)

---

```
1:  $r_0 \leftarrow b - Ax_0$      $z_0 \leftarrow Mr_0$   
2:  $\gamma_0 \leftarrow (z_0, r_0)$      $norm_0 \leftarrow \sqrt{(z_0, z_0)}$     ▷ 2 Merged Ops  
3: MPI_Allreduce on  $\gamma_0, norm_0$   
4:  $Z_0 \leftarrow Az_0$     ▷ Overlap Comm and Comp  
  
5: for  $j=1,2,\dots$ ,until convergence do  
6:   if  $i$  then  $\beta_j \leftarrow \gamma_{j-1}/\gamma_{j-2}$   
7:   else  $\beta_j \leftarrow 0.0$   
8:    $p_j \leftarrow z_{j-1} + \beta_j p_{j-1}$     ▷ 3 Merged Operations  
9:    $s_j \leftarrow Z_{j-1} + \beta_j s_{j-1}$      $\delta_j \leftarrow (p_j, s_j)$   
10:  MPI_Allreduce on  $\delta_j$   
11:   $S_j \leftarrow Ms_j$     ▷ Overlap Comm and Comp  
12:   $\alpha_j \leftarrow \gamma_{j-1}/\delta_j$      $x_j \leftarrow x_{j-1} + \alpha_j p_j$   
13:   $r_j \leftarrow r_{j-1} - \alpha_j s_j$     ▷ 4 Merged Operations  
14:   $z_j \leftarrow z_{j-1} - \alpha_j S_j$   
15:   $\gamma_j \leftarrow (r_j, z_j)$      $norm_j \leftarrow \sqrt{(z_j, z_j)}$   
16:  MPI_Allreduce on  $\gamma_j, norm_j$   
17:   $Z_j \leftarrow Az_j$     ▷ Overlap Comm and Comp
```

---

- Overlaps Matvec and PC each with non-blocking allreduce
- Algorithm from Gropp 2010



# Pipelined PCG (PIPECG)

---

```
1:  $r_0 \leftarrow b - Ax_0$      $u_0 \leftarrow Mr_0$      $w_0 \leftarrow Au_0$ 

2:  $\delta_0 \leftarrow (w_0, u_0)$                                  $\triangleright$  3 Merged Operations
3:  $\gamma_0 \leftarrow (r_0, u_0)$      $norm_0 \leftarrow \sqrt{(u_0, u_0)}$ 

4: MPI_Allreduce on  $\delta_0, \gamma_0, norm_0$                      $\triangleright$  Overlap Comm
5:  $m_0 \leftarrow Mw_0$      $n_0 \leftarrow Am_0$                  $\triangleright$  and Comp

6: for  $j=1,2,\dots$ ,until convergence do
7:   if  $j > 1$  then  $\beta_j \leftarrow \gamma_{j-1}/\gamma_{j-2}$ 
8:      $\alpha_j \leftarrow \gamma_{j-1}/(\delta_{j-1} - \beta_j/\alpha_{j-1}\gamma_{j-1})$ 
9:   else  $\beta_j \leftarrow 0.0$      $\alpha_j \leftarrow \gamma_{j-1}/\delta_{j-1}$ 

10:   $z_j \leftarrow n_{j-1} + \beta_j z_{j-1}$                          $\triangleright$  11 Merged Operations
11:   $q_j \leftarrow m_{j-1} + \beta_j q_{j-1}$      $p_j \leftarrow u_{j-1} + \beta_j p_{j-1}$ 
12:   $s_j \leftarrow w_{j-1} + \beta_j s_{j-1}$      $x_j \leftarrow x_{j-1} + \alpha_j p_{j-1}$ 
13:   $u_j \leftarrow u_{j-1} - \alpha_j q_{j-1}$      $w_j \leftarrow w_{j-1} - \alpha_j z_{j-1}$ 
14:   $r_j \leftarrow r_{j-1} - \alpha_j s_{j-1}$      $\delta_j \leftarrow (w_j, u_j)$ 
15:   $\gamma_j \leftarrow (r_j, u_j)$      $norm_j \leftarrow \sqrt{(u_j, u_j)}$ 

16:  MPI_Allreduce on  $\delta_j, \gamma_j, norm_j$                  $\triangleright$  Overlap Comm
17:   $m_j \leftarrow Mw_j$      $n_j \leftarrow Am_j$                  $\triangleright$  and Comp
```

---

- Overlaps Matvec and PC with single non-blocking allreduce
- Algorithm from Ghysels and Vanroose 2014



# 2-iteration Pipelined PCG (PIPE2CG)

---

```
1:  $r_0 \leftarrow b - Ax_0$      $z_0 \leftarrow Mr_0$      $w_0 \leftarrow Az_0$ 
2:  $\delta_0 \leftarrow (z_0, w_0)$      $\beta_0 \leftarrow (z_0, r_0)$      $norm_0 \leftarrow \sqrt{(z_0, z_0)}$     ▷ 3 Merged Operations

3: MPI_Allreduce on  $\delta_0, \beta_0, norm_0$ 
4:  $p_0 \leftarrow Mw_0$      $q_0 \leftarrow Ap_0$     ▷ Overlap Comm and Comp
5:  $c_0 \leftarrow Mq_0$      $d_0 \leftarrow Ac_0$ 

6: for  $j=1,3,\dots$ ,until convergence do
7:   if  $j > 1$  then
8:      $\rho_{j-1} \leftarrow 1/(1 - (\gamma_{j-1}\beta_{j-1})/(\gamma_{j-2}\beta_{j-2}\rho_{j-2}))$ 
9:      $\gamma_{j-1} \leftarrow \beta_{j-2}/\delta_{j-2}$ 
10:     $\phi \leftarrow [\rho_{j-1}, -\rho_{j-1}\gamma_{j-1}, (1 - \rho_{j-1})]$ 
11:     $\delta_{j-1} \leftarrow \phi_0\phi_0\lambda_8 - 2\phi_0\phi_1\lambda_2 + 2\phi_0\phi_2\lambda_3 + \phi_1\phi_1\lambda_4 - 2\phi_1\phi_2\lambda_5 + \phi_2\phi_2\lambda_9$ 
12:     $\beta_{j-1} \leftarrow \phi_0\phi_0\lambda_0 - 2\phi_0\phi_1\lambda_8 + 2\phi_0\phi_2\lambda_7 + \phi_1\phi_1\lambda_2 - 2\phi_1\phi_2\lambda_3 + \phi_2\phi_2\lambda_9$ 
13:     $\rho_j \leftarrow 1/(1 - (\gamma_j\beta_{j-1})/(\gamma_{j-1}\beta_{j-2}\rho_{j-1}))$ 
14:     $\gamma_j \leftarrow \beta_{j-1}/\delta_{j-1}$ 
15:  else  $\rho_j \leftarrow 1$      $\gamma_j \leftarrow \beta_{j-1}/\delta_{j-1}$ 
16:  VecPipelined_PIPE2CG()

17: MPI_Allreduce on  $\lambda_0$  to  $\lambda_9$     ▷ Overlap Comm
18:  $c_j \leftarrow Mq_j$      $d_j \leftarrow Ac_j$     ▷ and Comp
19:  $g_j \leftarrow Md_j$      $h_j \leftarrow Ag_j$ 
```

---



# PIPE2CG Vector Operations

```
1:  $\mu_j \leftarrow 1 - \rho_j$        $\mu_{j-1} \leftarrow 1 - \rho_{j-1}$ 
2: if  $j > 1$  then
3:    $x_{j-1} \leftarrow \rho_{j-1}(x_{j-2} + \gamma_{j-1}z_{j-2}) + \mu_{j-1}x_{j-3}$ 
4:    $r_{j-1} \leftarrow \rho_{j-1}(r_{j-2} - \gamma_{j-1}w_{j-2}) + \mu_{j-1}r_{j-3}$ 
5:    $z_{j-1} \leftarrow \rho_{j-1}(z_{j-2} - \gamma_{j-1}p_{j-2}) + \mu_{j-1}z_{j-3}$ 
6:    $w_{j-1} \leftarrow \rho_{j-1}(w_{j-2} - \gamma_{j-1}q_{j-2}) + \mu_{j-1}w_{j-3}$ 
7:    $p_{j-1} \leftarrow \rho_{j-1}(p_{j-2} - \gamma_{j-1}c_{j-2}) + \mu_{j-1}p_{j-3}$ 
8:    $q_{j-1} \leftarrow \rho_{j-1}(q_{j-2} - \gamma_{j-1}d_{j-2}) + \mu_{j-1}q_{j-3}$ 
9:    $c_{j-1} \leftarrow \rho_{j-1}(c_{j-2} - \gamma_{j-1}g_{j-1}) + \mu_{j-1}c_{j-3}$ 
10:   $d_{j-1} \leftarrow \rho_{j-1}(d_{j-2} - \gamma_{j-1}h_{j-1}) + \mu_{j-1}d_{j-3}$ 
```

▷ 8 Merged Ops

```
11:  $x_j \leftarrow \rho_j(x_{j-1} + \gamma_j z_{j-1}) + \mu_j x_{j-2}$ 
12:  $r_j \leftarrow \rho_j(r_{j-1} - \gamma_j w_{j-1}) + \mu_j r_{j-2}$ 
13:  $z_j \leftarrow \rho_j(z_{j-1} - \gamma_j p_{j-1}) + \mu_j z_{j-2}$ 
14:  $w_j \leftarrow \rho_j(w_{j-1} - \gamma_j q_{j-1}) + \mu_j w_{j-2}$ 
15:  $p_j \leftarrow \rho_j(p_{j-1} - \gamma_j c_{j-1}) + \mu_j p_{j-2}$ 
16:  $q_j \leftarrow \rho_j(q_{j-1} - \gamma_j d_{j-1}) + \mu_j q_{j-2}$ 
17:  $\lambda_0 \leftarrow (z_j, w_j)$      $\lambda_1 \leftarrow (z_j, q_j)$      $\lambda_2 \leftarrow (z_j, w_{j-1})$ 
18:  $\lambda_3 \leftarrow (p_j, q_j)$      $\lambda_4 \leftarrow (p_j, w_{j-1})$      $\lambda_5 \leftarrow (z_{j-1}, w_{j-1})$ 
19:  $\lambda_6 \leftarrow (z_j, r_j)$      $\lambda_7 \leftarrow (z_j, r_{j-1})$      $\lambda_8 \leftarrow (z_{j-1}, r_{j-1})$ 
20:  $\lambda_9 \leftarrow (z_j, z_j)$ 
21:  $\delta_j \leftarrow \lambda_0$      $\beta_j \leftarrow \lambda_6$      $norm_j \leftarrow \sqrt{\lambda_9}$ 
```

▷ 16 Merged  
Ops





# Scalable PCG Methods

Method	VecOps	Allr	Overlap	Init Costs
PCG	6	2 Allr	None	1 PC
SAPCG	7	1 Allr	None	1 PC, 1 Matvec
NBPCG	8	2 lallr	Matvec or PC	1 PC, 1 Matvec
PIPECG	11	1 lallr	Matvec, PC	2 PC, 2 Matvec
PIPE2CG*	24	1 lallr	2 Matvec, 2 PC	3 PC, 3 Matvec

Table: Differences between each PCG method

\*PIPE2CG computes two iterations of PCG each iteration

- Custom versions of solvers are implemented in PETSc
- Matrices are stored in MPIAIJ sparse compressed row format
- Primary preconditioner is block-Jacobi incomplete-Cholesky



# Matrix Experiments - Finite Element




Figure: Weak scaling tests for median iteration runtimes for matrices with about 220k nonzeros per process for 2-d (9 and 18 nonzeros per row) and 3-d (80 nonzeros per row) finite element matrices

# Allreduce and Vector Operations

## ■ Allreduce:

- Investigated time spent in `MPI_Test()` and `MPI_Wait()` compared to blocking allreduce for 27-point Poisson matrix with 4k rows per core
- Potential for 10-20% overall speedups by further optimizing non-blocking allreduce

## ■ Vector Operations:

- Investigated impact of merged vs separate vector operations for 27-point Poisson matrix with 13k rows per core
- Showed vector operation speedups of about 30% for PIPE2CG, 25% for PIPECG, and 10% for NBPCG

